

BayesAss Edition 3.0 User's Manual

Bruce Rannala

©2007 University of California Davis

Last updated 6 Mar 2023

1 Installation

The latest version of BA3 (version 3.0.5) can be downloaded [here](#). Unzip the archive by double clicking the downloaded file. A folder will be created in your current directory containing the source code and/or precompiled binary files for various computer operating systems. There is also an examples directory containing two example input files with samples from either 2 or 3 populations. If you have a compiler and are adventurous you can try compiling the source code (see below), otherwise refer to the instructions below to use a precompiled binary file for your specific computer operating system. There are two different executable files for each operating system. The executable **BA3SNP** is intended for use with single nucleotide polymorphism (SNP) data and allows a maximum of 4 alleles per locus, 30,000 loci, 100 populations and 3000 individuals. The executable **BA3MSAT** is intended for use with microsatellite (MSAT) data and allows a maximum of 500 alleles per locus, 500 loci, 100 populations and 5000 individuals. In the examples and discussion below the SNP program **BA3SNP** will often be used – the commands are the same for both **BA3SNP** and **BA3MSAT** – you should always use the appropriate executable file for your dataset (SNP or microsatellite).

1.1 Mac OS X

Using Homebrew package manager

The recommended way to install BA3 on Mac (or Linux) is using the homebrew package manager. Once you have homebrew installed on your machine you can install BA3 using the following commands (executed in the Terminal):

```
brew tap brannala/ba3
brew install ba3
```

The program will be installed and the commands **BA3SNP** and **BA3MSAT** will be available in the Terminal. You will want to download the documentation and example files (BA3-docs-examples.zip) from [here](#).

Using precompiled executables

Download the zip archived file with the latest version of the software [here](#). Unzip the archive by double clicking the downloaded file. A folder will be created in your current directory containing the program executable files **BA3SNP** and

BA3MSAT, example data files (in subfolder examples), and this manual (in subfolder docs). Note that there are two binaries available: “M” for users of newer M1 or M2 CPU Macs and “I” for users of older Intel CPU Macs. The Intel executable will run on M1/M2 Macs but this is through emulation and is not recommended.

1.2 Windows

Download the zip archived file with the latest version of the software here. Choose either the file `BA3Windows32.zip` or `BA3Windows64.zip` depending on whether you are running a 32 bit (older) or 64 bit (newer) version of Windows. Unzip the archive by double clicking the downloaded file. A folder will be created in your current directory containing the program executable files `BA3SNP.exe` and `BA3MSAT.exe`, example data files (in subfolder examples), and this manual (in subfolder docs). There are no longer any shared libraries that need to be installed for the Windows version of BA3.

1.3 Linux

Download the zip or tar.gz archived file with the latest version of the software here. Unzip the archive by double clicking the downloaded file. A folder will be created in your current directory containing the program executable files `BA3SNP` and `BA3MSAT`, example data files (in subfolder examples), and this manual (in subfolder docs).

1.3 Compiling the program

You do not need to compile the program if you have been successful following either steps 1.1 or 1.2 above. The source code for the program is found in the source tarball distribution file named `BA3-...tar.gz` where indicate the version numbers. The program uses routines from the gnu scientific library (gsl) and this library (and header files) must be installed prior to compiling. The gsl library can be found here. It is recommended that you install the gsl library using a package manager such as apt in Ubuntu linux or homebrew on Macs. If using a command line C++ compiler (e.g., g++, c++, etc), with gsl installed in the standard location, simply execute the following terminal commands in the directory that contains the source tarball:

```
tar -xvzf BA3-*.tar.gz
cd BA3-*
./make all
```

This will create the executable files `BA3SNP` and `BA3MSAT` in the current directory and typing `./BA3` at the command prompt will then execute the program. The procedures for compiling in Windows using the Cygwin package are essentially identical.

2 Running the program

The BA3 program is a command line program. If you are familiar with the unix terminal you will find it straightforward to use as it adheres to standard unix conventions for command line options, etc. If you have never used a terminal (command line) program you can find a beginners guide [here](#). Detailed instructions for running the program on the Mac OS X (or other unix-based) operating system are provided below.

2.1 Getting BA3 up and running on Mac OS X or Unix

2.1.1 Running BA3 installed with Homebrew

Download the examples and documentation `BA3-docs-examples.zip` and uncompress. If you uncompress the files on your desktop then in the Terminal type:

```
cd /Users/<login>/Desktop/BA3-docs-examples
```

to run a 3 population example dataset using the microsatellite version of the program type:

```
BA3MSAT -v examples/3pop.txt
```

You can also test the SNP version of the program since the example datasets only have 2 alleles per locus:

```
BA3SNP -v examples/3pop.txt
```

Note that in the examples that follow the prefix `./` is used before the program command name (e.g., `BA3SNP`), for example

```
./BA3SNP -v test.txt
```

This prefix represents the current working directory in Unix and should only be used when the executable files were manually installed and the program is being run from the directory containing the executable files. If you installed the program using Homebrew then the executables will already be in your path and so the prefix is not needed. You would instead execute the above command as

```
BA3SNP -v test.txt
```

If you incorrectly use the prefix Unix will complain that the executable file does not exist `./BA3SNP: No such file or directory. #####`

2.1.2 Running BA3 installed from binaries on Mac or Linux

To run the program, you will first need to start the terminal application which can be found in the Applications/Utilities folder on a Mac. A short tutorial on using the Mac OS X terminal can be found [here](#). The following description assumes that you have unzipped the BA3 distribution file on the Desktop. If you have placed it elsewhere you will need to change the commands to indicate the correct file path. Once you open terminal you will see a command line prompt. On a Mac, at the prompt type:

```
cd /Users/<login>/Desktop/BA3.*/binary/macosex
```

where is replaced with your account login name. On another unix computer you will specify the path to the binary that you unpacked (or compiled from source code). The unix command `cd` is short for “change directory” and the above command changes the current working directory from the user’s home directory (the default) to the directory where the BA3 program binary resides. To run the program using an example data file with 3 populations (contained in the subdirectory `examples`) type the following command:

```
./BA3SNP -v examples/3pop.txt
```

The prefix `./` means “current directory” and tells the operating system to look for the program file named `BA3SNP` in the current working directory. The program option `-v` specifies “verbose” output and causes BA3 to print out more detailed information to the screen when the program is running. ##### 2.1.3 Program screen output You should see output similar to the following:

```
BayesAss Edition 3.0.5 (BA3)
Released: 3/6/2023
Bruce Rannala
Department of Evolution and Ecology at UC Davis

Input file: examples/3pop.txt
Output file: BA3out.txt
Individuals: 400 Populations: 3 Loci: 7 Missing genotypes: 6

Locus:(Number of Alleles)
```

```
loc0:2 loc1:3 loc2:3 loc3:3 loc4:2 loc5:2 loc6:2
```

```
logP(M): -1618.34 logL(G): 0.00 logL: -1618.34 % done: (0.10) % accepted: (0.14, 0.83, 1.00)
logP(M): -1619.67 logL(G): 0.00 logL: -1619.67 % done: (0.20) % accepted: (0.15, 0.83, 1.00)
logP(M): -1621.13 logL(G): 0.00 logL: -1621.13 % done: (0.30) % accepted: (0.15, 0.83, 1.00)
logP(M): -1625.90 logL(G): 0.00 logL: -1625.90 % done: (0.40) % accepted: (0.15, 0.83, 1.00)
logP(M): -1622.89 logL(G): 0.00 logL: -1622.89 % done: (0.50) % accepted: (0.15, 0.83, 1.00)
logP(M): -1620.15 logL(G): 0.00 logL: -1620.15 % done: (0.60) % accepted: (0.15, 0.83, 1.00)
logP(M): -1619.98 logL(G): 0.00 logL: -1619.98 % done: (0.70) % accepted: (0.15, 0.83, 1.00)
logP(M): -1622.04 logL(G): 0.00 logL: -1622.04 % done: (0.80) % accepted: (0.15, 0.83, 1.00)
logP(M): -1617.66 logL(G): 0.00 logL: -1617.66 % done: (0.90) % accepted: (0.15, 0.83, 1.00)
logP(M): -1619.30 logL(G): 0.00 logL: -1619.30 % done: (1.00) % accepted: (0.15, 0.83, 1.00)
```

```
MCMC run completed. Output written to BA3out.txt
```

The program will create an output file in the current working directory when it has finished running. By default the output file is named `BA3out.txt`. You can double click on this file to open it with the Mac text editor and see the results.

2.2 Getting BA3 up and running on Windows

You will first need to run the Windows "Command Prompt" program which will open a console that you can use to run the `BA3SNP.exe` or `BA3MSAT.exe` programs. At the command prompt use the "cd" command to move to the directory where the files `BA3SNP.exe` and `BA3MSAT.exe` are found. For example,

```
C:\Users\bruce>cd Desktop\BA3Windows32
C:\Users\bruce\Desktop\BA3Windows32
```

To run the example file `3pop.txt` verbosely use the command:

```
C:\Users\bruce\Desktop\BA3Windows32>BA3MSAT.exe -v examples\3pop.txt
```

3 Data file format

The BA3 program uses an input file format that is identical to that of earlier BayesAss releases. The input file should be in a plain text format. DO NOT use a word processor such as Word to create the input file without explicitly converting it to a text file format before use. One possible approach is to input the data into a spreadsheet program such as Excel and then save the file as a "space-delimited text file." Another approach is to install one of the many available free text editors such as emacs or vi on your computer. Each line of the input file should have the following format

```
indivID popID locID allele1 allele2
```

where `indivID` is a unique identifier for the individual, `popID` is a unique identifier of the individual's source population, `locID` is a unique identifier for the locus, and `allele1` and `allele2` are the allele labels for each allele of the individual's genotype. The order of the alleles on the line is arbitrary. Missing alleles are represented using a 0. If there are n individuals and L loci there will be $n \times L$ lines in the input file. See the example data files distributed with the program. If you have genomic data in VCF file format you can use the ugnix tools to convert the data to BA3 format as described here.

4 Command line options

The BA3 program has about a dozen command line options that allow you to control the way the program runs and the level of detail in the output that it produces. The command line options are given after the program name and before the input file name. For example,

```
./BA3SNP -v -i=10000000 -o myout.txt myin.txt
```

executes the program for 1 million iterations using verbose output, writing the output to the file `myout.txt` and using the input file `myin.txt`. Some options such as the option specifying the number of iterations, `-i`, take parameter values while others such as `-v` do not. Parameter values should follow the option

specifier and may, or may not, be separated from the option specifier by a space. For example, the following are all equivalent ways to specify 1,000,000 iterations:

```
./BA3SNP -i1000000 myin.txt
./BA3SNP -i 1000000 myin.txt
./BA3SNP -i 1000000 myin.txt
```

Table 1 lists all the command line options with a brief description of their parameters and effects. Each option is described in detail in the remainder of this section. Following Unix conventions, each command line option has two possible forms, a short (one letter) form preceded by - and a longer, one word form preceded by --, for example the "verbose output" command can be specified on the command line as either `-v` or `--verbose`. The longer forms are available solely because some persons find them easier to remember.

Option	Values	Effect
-a -deltaA	$0 < \Delta_A \leq 1.0$	Mixing parameter for allele frequencies
-b -burnin	Positive integer	Number of iterations to discard as burnin
-f -deltaF	$0 < \Delta_F \leq 1.0$	Mixing parameter for inbreeding coefficients
-g -genotypes	None	Output genotypes and migrant ancestries
-i -iterations	Positive integer	Number of iterations for MCMC
-m -deltaM	$0 < \Delta_M \leq 1.0$	Mixing parameter for migration rates
-n -sampling	Positive integer	Interval between samples for MCMC
-o -output	String	Output file name
-s -seed	Positive integer	Seed for random number generator
-p -nolikelihood	None	Fix likelihood to 1 and generate priors
-t -trace	None	Create a trace file to monitor convergence
-u -settings	None	Output options and parameter settings
-v -verbose	None	Use verbose screen output

Table 1: Options available for BA3 program

4.1 Random number generator seed

The option `-s` (`-seed`) is used to specify a positive integer used to "seed" the random number generator algorithm. A deterministic algorithm is used to generate

pseudorandom numbers during the MCMC such that the sequence of random numbers is entirely determined by the starting seed. Thus, separate runs of the program started using same seed will produce exactly the same outcome. To test whether the program is converging it is important to carry out several independent runs initiated with different seeds. To start the program using 10456 as the random number seed use the following command:

```
./BA3SNP -s104656
```

If no seed is specified the default seed is 10.

4.2 MCMC iterations, burn-in and sampling interval

The command line option `-i` (`-iterations`) specifies the number of iterations for the Markov chain Monte Carlo (MCMC) analysis. By default the program uses 5,000,000 iterations. The number of iterations is an important factor in determining whether a MCMC analysis has converged (see below). In general, a greater number of iterations will be more likely to insure convergence but the run-time of the program also increases in proportion to the number of iterations. The value of the number of iterations should be a positive integer. For example,

```
.\BA3 -i10000000 test.txt
```

will execute the program using the data file `test.txt` and carry out 10 million iterations. The option `-b` (`-burnin`) is used to specify a positive integer that is the number of iterations of the MCMC that are discarded before sampling begins to obtain a sample of values that will be used to estimate parameters. Burn-in length is chosen such that the chain is likely to have reached the stationary distribution before sampling begins. The burn-in length must obviously be less than the total number of iterations. For example,

```
./BA3SNP -i10000000 -b1000000 test.txt
```

will run the MCMC for 10 million iterations, discarding the first 1 million iterations. In this case, 9 million iterations are available for sampling. The option `-n` (`-sampling`) is used to specify a positive integer that is the interval between samples. This interval must obviously be less than the number of iterations minus the burn-in, but will typically be much smaller, perhaps 100 or 1000. For example,

```
./BA3SNP -i10000000 -b1000000 -n1000 test.txt
```

will run the MCMC for 10 million iterations, discarding the first 1 million iterations and sampling every 1000 iterations from the remaining 9 million iterations, producing a sample of 9000 observations from the chain that will be used to estimate parameters.

4.3 MCMC mixing parameters

For continuous parameters such as migration rates, allele frequencies and inbreeding coefficients, the size of the proposed change to the parameter value at each iteration of the MCMC can be adjusted. These adjustments are used to fine-tune the acceptance rates for proposals (see discussion below). There are 3 mixing parameter adjustments: `-a` (`-deltaA`), `-f` (`-deltaF`) and `-m` (`-deltaM`) that adjust the proposal size for the allele frequencies, inbreeding coefficients and migration rates, respectively. Each mixing parameter should be a number between 0 and 1, with the size of the proposed move being proportional to the magnitude of this number.

4.4 Options for printing output

By default, the output produced by BA3 is written to a file named `BA3out.txt` that the program creates in the current working directory. An alternative name for the output file can be specified using the option `-o` (`-output`). For example,

```
./BA3SNP -o myout.txt test.txt
```

executes the program using the input file `test.txt` and writes the output to a file named `myout.txt`. The option `-t` (`--trace`) specifies whether a trace output file is created that lists all the parameter values at each iteration of the MCMC run. If this option is specified a file named `BA3trace.txt` is created in the current working directory. This file can be used to monitor convergence of the MCMC by plotting the profile of the likelihood and prior values, as well as those of various parameters, over time, using a program such as Tracer (see below). The option `-g` causes detailed information regarding the individual multilocus genotypes and posterior probabilities of migrant ancestries to be written to a file named `BA3indiv.txt` created in the current working directory. The option `-u` specifies that current values of command line options are printed at the beginning of the output file (recommended). Finally, the option `-v` (`--verbose`) specifies that detailed information about the input data (number of populations, number of loci, number of individuals, and so on) is written to the computer screen during the run and that likelihoods and parameter acceptance rates are written to the computer screen as the program runs (recommended). This detailed output can be used to adjust mixing parameters during initial trial runs (see below). It is also useful for checking that the input file is in the correct format and the data are being read correctly by the program.

5 Recommendations for running BA3

To generate correct results using BA3 it is important to adjust the mixing parameters, use a sufficient number of iterations, discard enough iterations as burn-in, and carry out several independent runs (started with different random number seeds), examining the trace files for evidence of convergence and mixing and looking for consistency of the estimates between independent runs. Here I

will outline a general strategy for achieving this. I illustrate the strategy using the example data file `3pop.txt`. I will begin with an explanation of the screen output generated using option `-v`.

5.1 Understanding BA3 screen output

Running the BA3 program using the command `./BA3SNP -v examples/3pop.txt` produces the following screen output:

```
BayesAss Edition 3.0.5 (BA3)
Released: 3/6/2023
Bruce Rannala
Department of Evolution and Ecology at UC Davis

Input file: examples/3pop.txt
Output file: BA3out.txt
Individuals: 400 Populations: 3 Loci: 7 Missing genotypes: 6

Locus:(Number of Alleles)

loc0:2 loc1:3 loc2:3 loc3:3 loc4:2 loc5:2 loc6:2
```

The first two lines of screen output (following the program title) specify the names of the input and output files. The next line prints the number of individuals (in this case, 400), the number of populations (in this case, 3), the number of loci (in this case, 7), and the total number of missing genotypes (in this case, 6). This is followed by a line specifying the number of alleles present at each locus. You should check that all these values agree with the expectations for your data. Discrepancies can indicate that there is a formatting error and the input file is not being read correctly. Once the MCMC begins running the current state of the chain will be printed to screen as follows:

```
logP(M): -1618.76 logL(G): -3087.23 logL: -4705.99 % done: [0.07]
% accepted: (0.31, 0.25, 0.66, 0.76, 0.60)
```

The first value `logP(M)` is the log probability of the current configuration of migrant ancestries among individuals, conditional on the current migration rates. The second value `logL(G)` is the log-likelihood of the genotype data given the migrant ancestries of individuals and the current population allele frequencies. The third value `logL` is the sum of these two terms. The value in brackets (or parentheses) after `done` is the percentage of the total iterations that have been completed. This proportion is displayed in square brackets if the chain is still in the burn-in phase, otherwise it is displayed in parentheses. The final output after `% accepted` is the acceptance rate for proposed changes to each of the 5 parameters from left to right:

1. migration rates

2. individual migrant ancestries
3. allele frequencies
4. inbreeding coefficients
5. missing genotypes

5.2 Adjustment of mixing parameters

The acceptance rates for proposed changes to parameters 1, 3 and 4 in the above list (migration rates, allele frequencies and inbreeding coefficients, respectively) can be adjusted by changing the values of the respective mixing parameters. If the acceptance rate is too high, the chain does not mix well, often proposing values very near the current value (which are accepted) and failing to adequately explore the state space. If the acceptance rate is too low the chain rarely accepts the proposed moves which are too different from the current value – this also causes poor mixing. Empirical analyses suggest that an acceptance rate between 20% and 60% is optimal. In the above example, the acceptance rate for proposed changes to migration rate is about 31% which is adequate. However, the acceptance rates for proposed changes to the allele frequencies and inbreeding coefficients are 66% and 76% respectively, which are both a bit high. One can decrease the acceptance rate by proposing larger moves (or increase the rate by proposing smaller ones). In this case, we want to decrease the acceptance rate so we will try increasing the proposal step size for the mixing parameters associated with proposed moves of both the allele frequencies and inbreeding coefficients. The default values of all the mixing parameters are 0.10. We will try increasing the proposal step length to 0.30 for both these proposals. Stop the program by typing Control-C in the terminal, then start it again using the following command options:

```
./BA3SNP -v -a0.30 -f0.30 examples/3pop.txt
```

The output from the MCMC run is now as follows:

```
logP(M): -1618.75 logL(G): -3142.41 logL: -4761.16 % done: [0.07]
% accepted: (0.31, 0.24, 0.31, 0.45, 0.60)
```

This is much better, but the acceptance rate for proposed changes to the inbreeding coefficients is still a bit high at 45%. We therefore again kill the program run using Control-C and try again with the following mixing parameters:

```
./BA3SNP -v -a0.30 -f0.50 examples/3pop.txt
```

The output from the MCMC run is now as follows:

```
logP(M): -1618.38 logL(G): -3115.25 logL: -4733.63 % done: [0.08]
% accepted: (0.31, 0.24, 0.32, 0.33, 0.60)
```

The acceptance rates now look okay so we will next try some longer runs with these values for the mixing parameters and create a trace file to examine convergence. Note that it may not always be possible to obtain acceptance rates in the recommended target range. If the likelihood surface is very flat, for example, as

may occur with weakly informative data, acceptance rates above 0.6 may occur even with a proposal step length of 1. In such cases, the mixing may still be satisfactory as indicated by the trace plot or other MCMC diagnostics. Also note that because the individual migrant ancestry and the missing genotypes are both discrete parameters the proposal step lengths are not adjustable and the maximum achievable acceptance rate is

$$\max P_{\text{accept}} = 2(1 - P_{\text{max}}),$$

where P_{max} is the largest posterior probability associated with any of the parameters. For example, if migrant ancestry state 0 (non-migrant) has posterior probability 0.95 then the maximum posterior acceptance rate is $2(1 - 0.95) = 0.10$. Thus, for some datasets the posterior acceptance rates for these parameters may be very small but the MCMC results are still reliable. Enter the following to initiate a longer run with the random seed 100, creating a trace file and printing the output to a file named `run1out.txt`:

```
./BA3SNP -v -a0.30 -f0.50 -t -s100 -i10000000 -b1000000 -n100 \
-o run1out.txt examples/3pop.txt
```

5.3 Interpreting the output file

The contents of the output file `run1out.txt` are as follows:

Input file: `examples/3pop.txt`

Individuals: 400 Populations: 3 Loci: 7

Locus:(Number of Alleles)

loc0:2 loc1:3 loc2:3 loc3:3 loc4:2 loc5:2 loc6:2

Population Index -> Population Label:

0->pop0 1->pop1 2->pop2

Migration Rates:

```
m[0][0]: 0.9718(0.0115) m[0][1]: 0.0130(0.0100) m[0][2]: 0.0152(0.0090)
m[1][0]: 0.0878(0.0141) m[1][1]: 0.7338(0.0396) m[1][2]: 0.1784(0.0407)
m[2][0]: 0.0870(0.0179) m[2][1]: 0.2047(0.0326) m[2][2]: 0.7083(0.0292)
```

Inbreeding Coefficients:

```
pop0 Fstat: 0.2552(0.0367)
pop1 Fstat: 0.0810(0.0588)
pop2 Fstat: 0.2698(0.0891)
```

Allele Frequencies:

```

pop0
loc0>>
2:0.806(0.029) 1:0.194(0.029)
loc1>>
3:0.309(0.033) 2:0.601(0.035) 1:0.089(0.020)
loc2>>
2:0.751(0.031) 3:0.170(0.027) 1:0.078(0.018)
loc3>>
2:0.384(0.035) 1:0.292(0.032) 3:0.324(0.032)
loc4>>
1:0.825(0.028) 2:0.175(0.028)
loc5>>
2:0.234(0.031) 1:0.766(0.031)
loc6>>
2:0.237(0.031) 1:0.763(0.031)

pop1
loc0>>
2:0.512(0.065) 1:0.488(0.065)
loc1>>
3:0.323(0.068) 2:0.445(0.070) 1:0.232(0.075)
loc2>>
2:0.222(0.062) 3:0.593(0.061) 1:0.186(0.054)
loc3>>
2:0.663(0.069) 1:0.325(0.071) 3:0.012(0.012)
loc4>>
1:0.266(0.071) 2:0.734(0.071)
loc5>>
2:0.750(0.049) 1:0.250(0.049)
loc6>>
2:0.868(0.058) 1:0.132(0.058)

pop2
loc0>>
2:0.479(0.077) 1:0.521(0.077)
loc1>>
3:0.451(0.083) 2:0.176(0.068) 1:0.374(0.080)
loc2>>
2:0.195(0.070) 3:0.433(0.077) 1:0.372(0.072)
loc3>>
2:0.465(0.069) 1:0.506(0.069) 3:0.029(0.017)
loc4>>
1:0.121(0.068) 2:0.879(0.068)
loc5>>
2:0.862(0.053) 1:0.138(0.053)

```

```
loc6>>
2:0.779(0.062) 1:0.221(0.062)
```

The first few lines of output summarize properties of the data. Next, there is a line that maps an integer index to each population label. This is done simply to allow the between population migration matrix to be printed more concisely. Next is the matrix of inferred (posterior mean) migration rates and the standard deviation of the marginal posterior distribution for each estimate. A rough 95% credible set can be constructed as $\text{mean} \pm 1.96 \times \text{sdev}$. Note that $m[i][j]$ is the fraction of individuals in population i that are migrants derived from population j (per generation). Next are the mean posterior estimates (and standard errors) of inbreeding coefficients and allele frequencies for each locus and population.

5.4 Diagnosing convergence

Two simple ways to examine convergence are:

- conduct multiple runs initialized with different seeds and compare the posterior mean parameter estimates for concordance.
- analyze the trace file for each run using the Tracer program (available [here](#)).

The trace file for the log-probability of the above run is plotted in Figure 1. The burn-in iterations are indicated in light grey, sample iterations in black. Two things should be observed. First, the log-probability initially increases steeply during the burn-in phase but then oscillates around a plateau – this is often (but not always) the case when a chain has converged. Second, the oscillations are quite regular – there are no persistent lows or highs (valleys or hills) in the plot – this is one indication that the chain is mixing well and effectively sampling from the posterior distribution with less autocorrelation between successive samples of the chain than would be the case if valleys and hills existed. The analysis of MCMC output is a generic problem and lots of programs are available (in R and other statistics packages) for analysing the results produced by BA3.




Figure 1: Trace file for log probability in a BA3 analysis of example input file 3pop.txt created using the Tracer program

5.5 Interpreting the individual ancestry output in BA3indiv.txt

If the command line option `-g` is used the file `BA3indiv.txt` is created in the current working directory at the end of the MCMC run. If we use the command

```
./BA3SNP -v -g -a0.30 -f0.50 -t -s100 -i10000000 -b1000000 -n100 \
-o run1out.txt examples/3pop.txt
```

this produces a `BA3indiv.txt` file containing an entry for each individual such as the following (I have only included a few individuals here for illustration):

```
Individual: ind0 Source Popln: 0
Genotypes>>
loc0:2/2 loc1:3/3 loc2:2/2 loc3:2/2 loc4:?? loc5:2/1 loc6:2/1
Migrant ancestry>>
[0,0]:0.931 [1,0]:0.000 [2,0]:0.000
[0,1]:0.000 [1,1]:0.002 [2,1]:0.001
[0,2]:0.000 [1,2]:0.035 [2,2]:0.031
```

This entry is for individual `ind0` sampled from population 0. The genotypes of the individual are listed first followed by the posterior probabilities of migrant ancestry. The notation `[i, j]` : indexes the population source `i` and generation `j` (0=nonmigrant, 1=1st generation migrant, 2=second generation migrant) of migrant ancestry. For example, `[0,0]` is the category of nonmigrants from population 0. If `j = 0` the only possible non-zero entry is the source population for the individual (in this example population 0). If `j > 0` the only possible non-zero entries are for populations other than the source population. In this case, the probability that the individual is a non-migrant is 0.931, the probability that it is a first-generation migrant from population 1 is 0.002, the probability that it is a second generation migrant from population 1 is 0.035, and so on. Below is another example individual:

```
Individual: ind11 Source Popln: 0
Genotypes>>
loc0:2/2 loc1:1/1 loc2:2/3 loc3:1/1 loc4:2/2 loc5:2/1 loc6:2/2
Migrant ancestry>>
[0,0]:0.175 [1,0]:0.000 [2,0]:0.000
[0,1]:0.000 [1,1]:0.242 [2,1]:0.483
[0,2]:0.000 [1,2]:0.031 [2,2]:0.069
```

Here the individual appears most likely to be a first generation migrant from population 2 (probability 0.483) although there is also non-negligible probability associated with the possibilities that the individual is either a non-migrant (probability 0.175) or a first-generation migrant from population 1 (probability 0.242).

5.6 The Priors

The prior on allele frequencies is uniform Dirichlet. So, with two alleles the prior means are $1/2$, with three alleles they are $1/3$ and so on. The prior distribution of the F statistic is uniform on the interval $(0, 1)$ with a mean of $1/2$. The prior on migration rates is uniform with the constraint that $m_{ii} \geq 2/3$ and $\sum_i \sum_{j \neq i} m_{ij} \leq 1/3$. The prior means for n populations are

$$\bar{m}_{ii} = \frac{1}{n} + \frac{2}{3} \left(\frac{n-1}{n} \right),$$

and

$$\bar{m}_{ij} = \frac{1}{3n} \text{ for } i \neq j.$$

Estimates of the prior variances can be obtained by running the program with option "p." It is suggested that users compare the posterior densities for their data with these prior densities to assess the change from prior to posterior which indicates how much information is contained in the dataset.